

Tennis-draft

May 9, 2019

```
[1]: from unityagents import UnityEnvironment
import numpy as np

env = UnityEnvironment(file_name="Tennis_Linux_NoVis/Tennis.x86_64")

# get the default brain
brain_name = env.brain_names[0]
brain = env.brains[brain_name]
# reset the environment
env_info = env.reset(train_mode=True)[brain_name]
# number of agents
num_agents = len(env_info.agents)
print('Number of agents:', num_agents)
# size of each action
# actions between -1 and 1
action_size = brain.vector_action_space_size
print('Size of each action:', action_size)
# examine the state space
states = env_info.vector_observations
state_size = states.shape[1]
print('There are {} agents. Each observes a state with length: {}'.
      →format(states.shape[0], state_size))
print('The state for the first agent looks like:', states[0])
```

INFO:unityagents:

'Academy' started successfully!

Unity Academy name: Academy

Number of Brains: 1

Number of External Brains : 1

Lesson number : 0

Reset Parameters :

Unity brain name: TennisBrain

Number of Visual Observations (per agent): 0

Vector Observation space type: continuous

Vector Observation space size (per agent): 8

Number of stacked Vector Observation: 3

Vector Action space type: continuous
Vector Action space size (per agent): 2
Vector Action descriptions: ,

Number of agents: 2

Size of each action: 2

There are 2 agents. Each observes a state with length: 24

The state for the first agent looks like: [0. 0. 0.

0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. -6.65278625 -1.5
-0. 0. 6.83172083 6. -0. 0.]

```
[2]: # main function that sets up environments
      # perform training loop

      from buffer import ReplayBuffer
      from maddpg import MADDPG
      import torch
      import numpy as np
      from tensorboardX import SummaryWriter
      import os
      from utilities import transpose_list, transpose_to_tensor
      from collections import deque

      device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

      # for saving gif
      import imageio

      def seeding(seed=1):
          np.random.seed(seed)
          torch.manual_seed(seed)

      def pre_process(entity, batchsize):
          processed_entity = []
          for j in range(3):
              list = []
              for i in range(batchsize):
                  b = entity[i][j]
                  list.append(b)
              c = torch.Tensor(list)
              processed_entity.append(c)
          return processed_entity

      seeding()
```

```

# number of parallel agents
# parallel_envs = 4

# number of training episodes.
# change this to higher number to experiment. say 30000.
number_of_episodes = 3000
episode_length = 1000
batchsize = 128
# how many episodes to save policy and gif
save_interval = 1000
t = 0

scores_deque = deque(maxlen=100)
scores = []

# amplitude of OU noise
# this slowly decreases to 0
noise = 2
noise_reduction = 0.9999

# how many episodes before update
episode_per_update = 2

log_path = os.getcwd()+"/log"
model_dir= os.getcwd()+"/model_dir"

os.makedirs(model_dir, exist_ok=True)

# torch.set_num_threads(parallel_envs)
# env = envs.make_parallel_env(parallel_envs)

# keep 5000 episodes worth of replay
buffer = ReplayBuffer(500000,batchsize,0)

# initialize policy and critic
maddpg = MADDPG(seed=100)
logger = SummaryWriter(log_dir=log_path)
agent0_reward = []
agent1_reward = []

# training loop
# show progressbar
import progressbar as pb
widget = ['episode: ', pb.Counter(), '/', str(number_of_episodes), ' ',
          pb.Percentage(), ' ', pb.ETA(), ' ', pb.Bar(marker=pb.
→RotatingMarker()), ' ' ]

```

```

timer = pb.ProgressBar(widgets=widget, maxval=number_of_episodes).start()

best_score = 0
max_score = 0

# use keep_awake to keep workspace from disconnecting
for episode in range(0, number_of_episodes):

    timer.update(episode)

    for agent in maddpg.maddpg_agent:
        agent.noise.reset()

    reward_this_episode = np.zeros(2)
    env_info = env.reset(train_mode=True)[brain_name] # reset the environment
    state = env_info.vector_observations # get the current state (for each
→agent)

    #for calculating rewards for this particular episode - addition of all time
→steps

    # save info or not
    save_info = ((episode % save_interval) < episode==number_of_episodes)
    frames = []
    tmax = 0

# if save_info:
# frames.append(env.render('rgb_array'))

    r0 = 0
    r1 = 0
    for episode_t in range(episode_length):

        t += 1

        # explore = only explore for a certain number of episodes
        # action input needs to be transposed
        actions = maddpg.act(torch.tensor(state,dtype=torch.float).to(device),
→noise=noise)
        noise *= noise_reduction

        action = torch.stack(actions).detach().numpy()

        # step forward one frame
        # next_obs, next_obs_full, rewards, dones, info = env.
→step(actions_for_env)

```

```

env_info = env.step(action)[brain_name]
next_state = env_info.vector_observations
reward = env_info.rewards
done = env_info.local_done

cat_state = np.concatenate((state[0],state[1]))
cat_next_state = np.concatenate((next_state[0],next_state[1]))

# add data to buffer
#transition = (obs, obs_full, actions_for_env, rewards, next_obs,
→next_obs_full, dones)
    transition = (state, cat_state, action, reward, next_state,
→cat_next_state, done)

    buffer.add(*transition)

    r0 += reward[0]
    r1 += reward[1]
    state = next_state

    # save gif frame
    if episode % 100 == 0 :
#         frames.append(env.render('rgb_array'))
        tmax+=1
    #if any of the agents are done break
    if len(buffer) > batchsize and episode % episode_per_update == 0:
        for a_i in range(2):
            samples = buffer.sample()
            maddpg.update(samples, a_i, logger)
            maddpg.update_targets() #soft update the target network towards the
→actual networks

        if np.any(done):
            break
    # update once after every episode_per_update

    agent0_reward.append(r0)
    agent1_reward.append(r1)
    r = max(r0,r1)

    scores.append(r)
    scores_deque.append(r)

    if episode % 100 == 0:
        avg_rewards = [np.mean(agent0_reward), np.mean(agent1_reward)]
        agent0_reward = []
        agent1_reward = []

```